

Root cause analysis via difference graph discovery from linear time-series data

Anouk Ruer¹ (✉), Timothée Loranchet¹, Daria Bystrova¹, and
Charles K. Assaad¹

Sorbonne Université, INSERM, Institut Pierre Louis d’Epidémiologie et de Santé
Publique, F75012, Paris, France {anouk.ruer}@iplesp.upmc.fr

Abstract. Root cause analysis aims to identify the mechanisms responsible for anomalies in complex dynamical systems. In this paper, we study root cause analysis in linear time-series through the lens of difference graph discovery. We focus on effect-defying root causes, corresponding to variables whose causal coefficients change between a normal and an anomalous regime. We formalize this problem using linear discrete-time dynamic structural causal models and adapt several methods originally introduced for discovering difference graphs between two populations to the time-series setting, where the two populations are replaced by a normal and an anomalous regime. We first evaluate the proposed approaches on simulated data, and then demonstrate their practical relevance on real-world datasets from IT monitoring and intensive care monitoring. Our results show how difference graph discovery can help localize causal mechanisms responsible for anomalous behavior.

Keywords: difference graphs · root cause analysis · time-series.

1 Introduction

Root cause analysis is a central task in many scientific and industrial domains, where the goal is to identify the mechanisms responsible for observed changes in a system. In climate science, for instance, understanding the drivers of extreme events or long-term shifts is crucial for prediction and policy-making [9]. In IT monitoring, root cause analysis supports the diagnosis of system failures and performance degradation, enabling faster and more reliable interventions [4]. Similarly, in healthcare, identifying the causes of adverse events or changes in patient trajectories can have important clinical implications, for example in intensive care or nephrology settings [14,11,17,1].

In parallel, there has been growing interest in methods that learn differences between two environments [28,10,18,8,7]. These approaches, often referred to as difference graph discovery methods, aim to identify changes in dependency or causal structures across datasets. Although they have been studied in several settings, their potential for root cause analysis remains largely unexplored.

In this paper, we investigate how difference graph discovery can be leveraged for root cause analysis in a linear setting. A key limitation of existing methods is

that they are mostly designed for static data, whereas many real-world systems evolve over time. We therefore extend difference graph discovery methods to the time-series setting, allowing them to capture dynamic changes in causal or dependency structures between a normal and an anomalous regime. More specifically, we extend DCI [28] and LDiffPC [8,7], and provide theoretical guarantees for their correctness in the time-series setting.

We highlight that these methods are naturally suited to detecting certain types of root causes, in particular effect-defying root causes, corresponding to variables whose causal coefficients change between a normal and an anomalous regime [4,1]. We then empirically evaluate the proposed extensions on both simulated and real-world datasets. For the real-world evaluation, we consider a dataset from IT monitoring as well as a dataset from intensive care monitoring. These experiments assess the practical relevance of difference graph discovery for root cause analysis and illustrate its potential across different application domains.

The remainder of the paper is organized as follows. Section 2 introduces the necessary preliminaries. Section 3 reviews related work. Section 4 presents the time-series extensions of the considered difference graph discovery algorithms. Section 5 discusses how these algorithms can detect effect-defying root causes. Section 6 reports experiments on simulated data and apply the proposed methods to IT monitoring and intensive care monitoring data, respectively. Finally, Section 7 discusses the results and concludes the paper. All proofs are provided in Appendix. The implementation of the extended time-series difference graph methods, together with the code used for the experiments is available at <https://github.com/CIPHOD/pyCIPHOD/tree/main/reproducibility/ecmlpkddcaesar2026>.

2 Background

We consider multivariate time-series $\mathbb{V}$ evolving according to an unknown linear discrete-time dynamic structure causal model (DT-DSCM) [22,4], denoted by $\mathcal{M}$, such that $\forall Y_t \in \mathbb{V}, Y_t = \sum_{X_{t-\ell} \in \mathbb{V}} \alpha_{X_{t-\ell}, Y_t} X_{t-\ell} + \varepsilon_{Y_t}$, where $\alpha_{X_{t-\ell}, Y_t} \neq 0$ only if $X_{t-\ell}$ is a direct cause of Y_t and ε_{Y_t} is an exogenous noise term. Variables with a temporal subscript, such as Y_t or $X_{t-\ell}$, denote time-indexed endogenous variables. When referring to an entire time-series, we omit the temporal subscript and write Y or X . The DT-DSCM governs the temporal and causal dependencies among the variables in $\mathbb{V}$. The model respects temporal ordering, meaning that a variable X_t cannot cause $Y_{t-\ell}$ for any $\ell > 0$, while allowing instantaneous relations, so that X_t can directly cause Y_t . In this paper we assume that there is no unmeasured confounding, an assumption known as causal sufficiency.

Assumption 1 (Causal sufficiency) *All noise terms are mutually independent and each noise term can affect only one observed variable.*

We consider the presence of two regimes $\mathcal{N}$ and $\bar{\mathcal{N}}$, corresponding to two different conditions of the system, namely, $\mathcal{N}$ represents the normal regime and




Fig. 1: FT-DAG with a regime change at time t and $\ell_{max} = 1$. The vertices highlighted in orange are affected by a change in its causal mechanism; in this example, the time-series Y is an effect-defying root cause. We consider two settings depending on how the causal mechanism of Y changes: (i) Setting R, where only the red edges are changed, and (ii) Setting RB, where both the red and blue edges are changed. Note that the causal mechanism can also be changed through a lagged mechanism, such as $Y_{t-1} \rightarrow Z_t$, where Z would be the root cause.

$\tilde{\mathcal{N}}$ represents the anomalous regime. Throughout the paper, we assume that the observations have been assigned to either the normal or the anomalous regime before applying the difference graph discovery methods. This assumption is standard in both the difference graph discovery literature, where the generating regime is assumed to be known, and the root cause analysis literature, where anomalous observations are assumed to be identified beforehand, as is the case for all methods reviewed in Section 3. In the following, we use superscripts to indicate the regime associated with a given object. For example, $\mathcal{M}^{\mathcal{N}}$, $\alpha_{X_{t-\ell}, Y_t}^{\tilde{\mathcal{N}}}$, and $\varepsilon_{Y_t}^{\tilde{\mathcal{N}}}$ denote, respectively, the DT-DSCM in the anomalous regime, the causal coefficient from $X_{t-\ell}$ to Y_t in the anomalous regime, and the exogenous noise of Y_t in the anomalous regime. Within each regime the underlying DT-DSCM is assumed to be causally stationary.

Assumption 2 (Causal stationarity) *Within each regime, for each variable Y_t , the coefficients $\alpha_{X_{t-\ell}, Y_t}$ do not depend on t and the noise ε_{Y_t} is identically distributed over time.*

In addition, we assume the existence of a finite maximal lag between causes and their effects in the system, denoted by ℓ_{max} , which is much smaller than the total number of observed time points. In practice, we assume that this maximal lag is provided by domain knowledge. More generally, the domain expert may specify an upper bound on the true maximal lag, so that the chosen ℓ_{max} is greater than or equal to the actual maximal lag of the system.

Within each regime, the DT-DSCM induces a full-time directed acyclic graph (FT-DAG), denoted by $\mathcal{G}^{\mathcal{N}} = (\mathbb{V}, \mathbb{E}^{\mathcal{N}})$ and $\mathcal{G}^{\tilde{\mathcal{N}}} = (\mathbb{V}, \mathbb{E}^{\tilde{\mathcal{N}}})$, where vertices correspond to time-indexed variables and edges encode direct causal relationships both within a time step and across successive time steps. In this paper, we use several standard graphical notions. For instance, a vertex $X_{t-\ell}$ is a parent of

Y_t in $\mathcal{G}$ if $X_{t-\ell} \rightarrow Y_t$ in $\mathcal{G}$, equivalently if $\alpha_{X_{t-\ell}, Y_t} \neq 0$. The set of parents of Y_t is denoted by $\text{Pa}(Y_t, \mathcal{G})$. Two vertices are adjacent if they are connected by an edge in $\mathcal{G}$. The skeleton of $\mathcal{G}$, denoted by $\text{Skel}(\mathcal{G})$, is the undirected graph obtained from $\mathcal{G}$ by replacing every directed edge with an undirected edge while preserving the same set of adjacencies.

For illustration, Figure 1 presents FT-DAGs corresponding to a normal and an anomalous regime. In the normal regime, the system is governed by a fixed stationary DT-DSCM, whereas in the anomalous regime it is governed by another stationary DT-DSCM in which changes, relative to the normal regime, occur in the direct effects on a specific variable Y_t . Two settings are considered in this figure. In the first one, referred to as *Setting R*, only the red edge from X_t to Y_t is modified across regimes. In the second one, referred to as *Setting RB*, both the red and blue edges, from X_t to Y_t and from Z_t to Y_t , are modified. These settings will serve as running examples throughout the paper to clarify the behavior of the proposed methods.

Under Assumption 2, the FT-DAG of each regime is fully characterized by the relationships between variables within a window of size $\ell_{\max} + 1$, as it repeats this structure throughout the regime. Henceforth, $\mathbb{V}$ refers solely to variables within this window; see Appendix for more details.

The main objective of this paper is to identify root causes. To relate observed anomalies to changes in the causal system, we make the following assumption.

Assumption 3 (Anomalies) *All anomalies observed in the anomalous regime are induced by changes in the underlying DT-DSCM between the normal regime $\mathcal{N}$ and the anomalous regime $\tilde{\mathcal{N}}$.*

Under the above assumption, root causes are defined as variables whose changes trigger and propagate anomalies throughout the system. At the finest level, such changes occur at the micro, or time-indexed, level of the system, for instance through a change in the mechanism of a variable Y_t . However, under stationarity, a change affecting Y_t at the beginning of the anomalous regime may propagate throughout the rest of the regime. Therefore, for simplicity and practical interpretability, we define root causes at the macro, or time-series, level. That is, we say that the time-series Y is a root cause, rather than referring to a specific time point Y_t . This is also consistent with many applications, where domain experts are primarily interested in identifying which process is causing the malfunctioning, rather than the exact time at which the malfunctioning first appeared.

In DT-DSCMs, two types of root causes are commonly considered [1]. A macro vertex Y is called an *effect-defying root cause* if there exist $Y_t \in Y$ and $X_{t-\ell} \in \text{Pa}(Y_t, \mathcal{G}^{\mathcal{N}})$ such that the corresponding causal coefficient changes across regimes, namely $\alpha_{X_{t-\ell}, Y_t}^{\mathcal{N}} \neq \alpha_{X_{t-\ell}, Y_t}^{\tilde{\mathcal{N}}}$. Similarly, Y is called a *noise-defying root cause* if there exists $Y_t \in Y$ such that the noise distribution of Y_t differs between regimes, that is, $\varepsilon_{Y_t}^{\mathcal{N}} \neq_d \varepsilon_{Y_t}^{\tilde{\mathcal{N}}}$.

In this paper, we mainly focus on effect-defying root causes. Difference graphs provide a natural tool for this task, as they aim to identify direct relationships that differ between variables across two distributions or environments. While




Fig. 2: FT-DFGs corresponding to the FT-DAG in Figure 1. The left panel shows the FT-DFG for Setting R, where only the red edges in Figure 1 are changed. The right panel shows the FT-DFG for Setting RB, where both the red and blue edges in Figure 1 are changed.

difference graph methods have received increasing attention in the independent and identical distributed (i.i.d.) setting, mostly to characterize differences between populations [28,10,18,8,2] rather than for root cause analysis, their role in time-series root cause analysis remains largely unexplored. Here, we investigate how such methods can be used to detect effect-defying root causes in linear dynamical systems by comparing normal and anomalous regimes. In the following, we introduce difference graphs adapted to time-series and linear DT-DSCMs.

Definition 1 (Full-time Difference Graph). Consider two DT-DSCMs $\mathcal{M}^{\mathcal{N}}$ and $\mathcal{M}^{\tilde{\mathcal{N}}}$. A Full-Time Difference Graph (FT-DFG) $\mathcal{D} = (\mathbb{V}, \mathbb{E}^{|\mathcal{N}-\tilde{\mathcal{N}}|})$ between $\mathcal{M}^{\mathcal{N}}$ and $\mathcal{M}^{\tilde{\mathcal{N}}}$, is a directed graph such that the set of vertices is identical to the set of endogeneous variables in $\mathcal{M}^{\mathcal{N}}$, $\mathcal{M}^{\tilde{\mathcal{N}}}$ and the set of edges is defined as: $\mathbb{E}^{|\mathcal{N}-\tilde{\mathcal{N}}|} = \{X_{t-\ell} \rightarrow Y_t | \forall (X_{t-\ell}, Y_t) \in \mathbb{V}^2 \text{ such that } \alpha_{X_{t-\ell}, Y_t}^{\mathcal{N}} \neq \alpha_{X_{t-\ell}, Y_t}^{\tilde{\mathcal{N}}}\}$.

In general, FT-DFGs may contain cycles [2]. However, to simplify the problem, we assume in this paper that the FT-DFG is acyclic. This assumption is commonly adopted by algorithms [28,10,18,8] that aim to fully or partially recover FT-DFGs from data.

Assumption 4 (Common topological ordering) The graphs $\mathcal{G}^{\mathcal{N}}$ and $\mathcal{G}^{\tilde{\mathcal{N}}}$ share a common topological ordering. That is, there exists a mapping $\pi : \mathbb{V} \rightarrow \{1, \dots, |\mathbb{V}|\}$ such that π is a topological ordering of both $\mathcal{G}^{\mathcal{N}}$ and $\mathcal{G}^{\tilde{\mathcal{N}}}$, i.e., for every directed edge $X_t \rightarrow Y_t$ in either graph, $\pi(X_t) < \pi(Y_t)$.

Figure 2 illustrates the FT-DFGs associated with the two settings, R and RB, from Figure 1.

3 Related Work

Existing root-cause analysis methods either rely on a predefined causal graph [5,16,4,1] or first estimate causal relations from data [27,25,24,3,23]. Among the latter, MicroCause [20] combines PCMCI [25] with a random-walk procedure, RCD [12] uses a regime indicator, and T-RCA [29] combines a learned graph with anomaly timing. There exists also approaches that focus on explaining one-point anomalies or outliers [6,21,26] but we consider those to be beyond the scope of this

paper. Here we focus on collective anomalies and assume having access to samples from normal and anomalous regimes.

Our work builds on difference-graph discovery, which directly estimates how causal mechanisms vary across environments. In particular, we focus on methods that use equality tests to infer the FT-DFG.

- LDiffPC [8,7] is an algorithm for discovering a partially oriented graph in linear non-dynamic structural causal models. It recovers the skeleton of the graph by testing equality of regression coefficients across environments and then orients edges using collider detection and Meek [19] propagation rules.
- DCI [28] is also designed to discover a partially oriented graph in linear non-dynamic structural causal models. It recovers the skeleton using the same procedure as LDiffPC. However, its orientation step differs: instead of relying on separation sets and orientation rules, DCI orients edges by exploiting changes in residual variances across environments.

There exists also other approaches that do not use equality tests. For instance, MBGH [18] relies on covariance differences, while iSCAN [10] identifies changed mechanisms in nonlinear additive-noise models through distributional invariance. More details on related works is provided in Appendix.

4 Difference graph discovery: from static to time-series

In this section, we adapt existing difference graph discovery methods to the time-series setting. The methods considered here do not return a fully oriented FT-DFG, but rather a partially oriented graph. We denote the output of such methods by $\hat{\mathcal{D}} = (\mathbb{V}, \hat{\mathbb{E}}, \hat{\mathbb{E}}^-)$, where $\mathbb{V}$ is the set of vertices, $\hat{\mathbb{E}}$ is the set of directed edges, and $\hat{\mathbb{E}}^-$ is the set of undirected edges.

In the following, σ_{Y_t} denotes the standard deviation of the noise term ε_{Y_t} . For any conditioning set $\mathbb{S}$, $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}}$ denotes the coefficient of $X_{t-\ell}$, and $\sigma_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}}$ denotes the residual standard deviation, in the regression of Y_t on $\{X_{t-\ell}\} \cup \mathbb{S}$ in the normal regime. Similarly, $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{A}}$ and $\sigma_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{A}}$ denote the corresponding coefficient and residual standard deviation in the anomalous regime.

4.1 The tsLDiffPC algorithm

We begin with the extension of LDiffPC, which we denote by tsLDiffPC. tsLDiffPC proceeds in two steps. First, it estimates the FT-DFG skeleton by testing equality of regression coefficients across regimes. More specifically, tsLDiffPC starts from a complete undirected graph over $\mathbb{V}$ and iteratively removes edges between pairs of vertices $X_{t-\ell}$ and Y_t whenever it finds a conditioning set $\mathbb{S} \subseteq \mathbb{V} \setminus \{X_{t-\ell}, Y_t\}$ such that $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}} = \beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{A}}$. The size of the conditioning set is increased progressively, following the PC strategy [27]. For each removed edge, the corresponding separating set $\mathbb{S}$ is stored. Second, tsLDiffPC starts the orientation phase by orienting lagged edges according to the temporal order. It




Fig. 3: FT-DFGs recovered by the considered algorithms. The left panel shows the output of tsLDiffPC in Setting R. The middle panel shows the output of tsLDiffPC² and tsDCIPC in Setting R. Under the variance-invariance condition required by DCI, tsDCI recovers the same graph. The right panel shows the output of tsLDiffPC, tsLDiffPC², and tsDCIPC in Setting RB; under the same condition, tsDCI recovers an identical graph.

then orients instantaneous edges using the standard LDiffPC collider-detection rule: for any triple $X_{t-\ell} - Z_t - Y_t$ such that $X_{t-\ell}$ and Y_t are not adjacent, if Z_t does not belong to the separating set stored for $X_{t-\ell}$ and Y_t , the triple is oriented as $X_{t-\ell} \rightarrow Z_t \leftarrow Y_t$. Finally, tsLDiffPC further orients edges by iteratively applying Meek's orientation rules until no additional orientation is possible. The pseudocode is given in Appendix.

To ensure correct detection, tsDiffPC relies on the following two assumptions.

Assumption 5 (Diff-adjacency-faithfulness) Let $\mathcal{D} = (\mathbb{V}, \mathbb{E})$ denote an FT-DFG. If $X_{t-\ell} \rightarrow Y_t$ in $\mathbb{E}$ then for any subset $\mathbb{S} \subseteq \mathbb{V} \setminus \{X_{t-\ell}, Y_t\}$, $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}} \neq \beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}}$.

Assumption 6 (Diff-orientation-faithfulness) Let $\mathcal{D} = (\mathbb{V}, \mathbb{E})$ denote an FT-DFG. Let $(X_{t-\ell}, Z_t, Y_t)$ be a triplet of vertices such that $X_{t-\ell}$ and Z_t are adjacent, Z_t and Y_t are adjacent, and $X_{t-\ell}$ and Y_t are not adjacent. The following holds:

- If $X_{t-\ell} \rightarrow Z_t \leftarrow Y_t$ in $\mathcal{D}$ then for any subset $\mathbb{S} \subseteq \mathbb{V} \setminus \{X_{t-\ell}, Y_t\}$ that contains Z_t , $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}} \neq \beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}}$.
- Otherwise, for any subset $\mathbb{S} \subseteq \mathbb{V} \setminus \{X_{t-\ell}, Y_t\}$ that does not contain Z_t , $\beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}} \neq \beta_{X_{t-\ell}, Y_t | \mathbb{S}}^{\mathcal{N}}$.

The two assumptions introduced above differ from those originally used in the work introducing LDiffPC. In [7], the analysis relies on the stronger assumption of diff-faithfulness, which establishes an exact correspondence between the equality of regression coefficients and diff-separation (a tool similar to d-separation [22] but adapted to difference graphs). The analysis also requires an additional assumption ensuring the symmetry of regression invariance. It was shown that diff-adjacency-faithfulness and diff-orientation-faithfulness are jointly implied by diff-faithfulness and the symmetry of regression invariance [7]. Here, we argue that diff-adjacency-faithfulness and diff-orientation-faithfulness are sufficient on their own to establish the correctness of the tsLDiffPC algorithm, as well as its

extension to the time series setting. Further details on how Assumptions 5 and 6 relate to the assumptions used in [7] are provided in Appendix.

Theorem 1 establishes the correctness of tsLDiffPC under diff-adjacency-faithfulness and the diff-orientation-faithfulness introduced above.

Theorem 1. *Let $\mathcal{D} = (\mathbb{V}, \mathbb{E})$ be the true FT-DFG and $\hat{\mathcal{D}} = (\mathbb{V}, \hat{\mathbb{E}}, \hat{\mathbb{E}}^-)$ be the output of tsLDiffPC. If Assumptions 1,2,3,4, 5 and 6 are satisfied and we are given perfect conditional equality information about all pairs of variables then $\text{Skel}(\hat{\mathcal{D}}) = \text{Skel}(\mathcal{D})$ and $\hat{\mathbb{E}} \subseteq \mathbb{E}$.*

Example 1. In Setting R, only a single edge entering Y_t is modified across regimes. Consequently, the corresponding FT-DFG contains only one changed edge, $X_t \rightarrow Y_t$, as shown in Figure 2 (left). While LDiffPC correctly identifies the adjacency, it lacks sufficient information to orient the edge, as neither temporal ordering, collider detection, nor Meek’s orientation rules apply. As a result, the edge remains unoriented in the output, shown in Figure 3 (left). In contrast, in Setting RB, two incoming edges into Y_t are modified across regimes. The resulting FT-DFG contains an unshielded collider structure of the form $X_t \rightarrow Y_t \leftarrow Z_t$, where Y_t is the common endpoint of the changed edges. In this case, LDiffPC exploits its collider orientation rule to orient both edges toward Y_t (Figure 3 right).

4.2 The tsDCI algorithm

The second algorithm, denoted tsDCI is a time-series extension of DCI. The tsDCI algorithm proceeds in two steps. First, it estimates the FT-DFG skeleton using the same procedure as tsLDiffPC. Lagged edges are oriented directly according to temporal order, while instantaneous edges are oriented by testing the invariance of residual variances across regimes. More precisely, tsDCI searches for a conditioning set $\mathbb{S}$ such that $\sigma_{X_t|\mathbb{S}}^{\mathcal{N}} = \sigma_{X_t|\mathbb{S}}^{\tilde{\mathcal{N}}}$. If such a set $\mathbb{S}$ exists and $Y_t \notin \mathbb{S}$, the edge is oriented as $X_t \rightarrow Y_t$; if $Y_t \in \mathbb{S}$, it is oriented as $Y_t \rightarrow X_t$. The procedure is repeated until no further orientation is possible. The pseudocode is given in Appendix.

To ensure correctness, tsDCI requires, similarly to tsLDiffPC, the diff-adjacency-faithfulness assumption. But for orientation, it relies on the following assumption on residual variances, which is the temporal analogue of Assumption 4.2 in [28].

Assumption 7 (Var-orientation-faithfulness) *(temporal analogue of Assumption 4.2 [28]) For any node $X_{t-\ell}, Y_t \in \mathbb{V}$ it holds that:*

- If $\alpha_{X_{t-\ell}, Y_t}^{\mathcal{N}} \neq \alpha_{X_{t-\ell}, Y_t}^{\tilde{\mathcal{N}}}$ then for every $\mathbb{S} \subseteq \mathbb{V} \setminus \{Y_t\}$, $\sigma_{Y_t|\mathbb{S}}^{\mathcal{N}} \neq \sigma_{Y_t|\mathbb{S}}^{\tilde{\mathcal{N}}}$ and $\sigma_{X_{t-\ell}|\mathbb{S} \cup \{Y_t\}}^{\mathcal{N}} \neq \sigma_{X_{t-\ell}|\mathbb{S} \cup \{Y_t\}}^{\tilde{\mathcal{N}}}$.
- If $\sigma_{Y_t}^{\mathcal{N}} \neq \sigma_{Y_t}^{\tilde{\mathcal{N}}}$, then $\sigma_{Y_t|\mathbb{S}}^{\mathcal{N}} \neq \sigma_{Y_t|\mathbb{S}}^{\tilde{\mathcal{N}}}$ for every $\mathbb{S} \subseteq \mathbb{V} \setminus \{Y_t\}$.

Theorem 2 establishes the correctness of tsDCI under diff-adjacency-faithfulness and the residual-variance assumption introduced above.

Theorem 2. *Let $\mathcal{D} = (\mathbb{V}, \mathbb{E})$ be the true FT-DFG and $\hat{\mathcal{D}} = (\mathbb{V}, \hat{\mathbb{E}}, \hat{\mathbb{E}}^-)$ be the output of tsDCI. If Assumptions 1,2,3,4,5 and 7 are satisfied and we are given perfect conditional equality information about all pairs of variables then $\text{Skel}(\hat{\mathcal{D}}) = \text{Skel}(\mathcal{D})$ and $\hat{\mathbb{E}} \subseteq \mathbb{E}$.*

Example 2. Similarly to tsLDiffPC, tsDCI recovers the adjacencies in both settings R and RB. However, unlike tsLDiffPC, tsDCI can orient certain isolated changed edges. For example in Setting R, if $\sigma_{Y_t}^{\mathcal{N}} = \sigma_{Y_t}^{\bar{\mathcal{N}}}$, then Assumption 7 guarantees the existence of a conditioning set witnessing residual-variance invariance for Y_t . Consequently, tsDCI orients the edge toward Y_t , thereby recovering the true FT-DFG as shown in the middle panel of Figure 3.

4.3 The tsLDiffPC² and the tsDCIPC algorithms

We also consider two augmented variants, tsLDiffPC² and tsDCIPC. The main idea behind these algorithms is to first apply a causal discovery method to the normal data, and then leverage the inferred orientations by projecting them onto the FT-DFG. To enable this step, we require the following additional assumption.

Assumption 8 (Anomalous subgraph) *Let $\mathcal{G}^{\mathcal{N}} = (\mathbb{V}, \mathbb{E}^{\mathcal{N}})$ and $\mathcal{G}^{\bar{\mathcal{N}}} = (\mathbb{V}, \mathbb{E}^{\bar{\mathcal{N}}})$ denote the FT-DAGs associated with the normal and anomalous regimes, respectively. It is assumed that $\mathbb{E}^{\bar{\mathcal{N}}} \subseteq \mathbb{E}^{\mathcal{N}}$.*

This assumption is common in the root cause analysis literature and is notably adopted by EasyRCA and SGRCA. An alternative would be to assume that the normal graph is a subgraph of the anomalous graph and to apply tPC to the anomalous regime instead. We focus on the former setting because anomalies are more likely to modify existing mechanisms than to create entirely new ones, and because the normal regime typically contains more observations, leading to more reliable causal discovery.

The tsLDiffPC² and tsDCIPC algorithms first run tsLDiffPC and tsDCI, respectively, to obtain a partially oriented FT-DFG. Then apply tPC [27] to the normal regime and estimate a partially oriented FT-DAG. As in the standard PC algorithm [27], tPC removes edges by testing conditional independences and subsequently orients the remaining edges using temporal order, the collider-detection rule, and Meek rules. Finally, each undirected edge in $\hat{\mathcal{D}}$ is oriented whenever tPC provides an orientation in $\hat{\mathcal{G}}^{\mathcal{N}}$. Since both methods rely on tPC, they additionally require the assumptions underlying tPC.

Assumption 9 (Faithfulness in the normal regime) *Let $\mathcal{M}^{\mathcal{N}}$ be the DT-DSCM corresponding to the normal regime, with causal graph $\mathcal{G}^{\mathcal{N}}$ and distribution $\mathbb{P}^{\mathcal{N}}$. All conditional independences in $\mathbb{P}^{\mathcal{N}}$ are implied by the causal Markov property associated with $\mathcal{G}^{\mathcal{N}}$, which states that each variable is independent of its non-descendants given its parents.*

Table 1 summarizes all assumptions required by each of the algorithms introduced in this paper. The soundness of tsLDiffPC² and tsDCIPC trivially follows from the soundness of tsLDiffPC, tsDCI and tPC.

Assumption	tsLDiffPC	tsLDiffPC ²	tsDCI	tsDCIPC
Linearity	✓	✓	✓	✓
Sufficiency (As. 1)	✓	✓	✓	✓
Stationarity (As. 2)	✓	✓	✓	✓
Anomalies (As. 3)	✓	✓	✓	✓
Order (As. 4)	✓	✓	✓	✓
D. a. faithful. (As. 5)	✓	✓	✓	✓
D. o. faithful. (Ass. 6)	✓	✓		
D. v. o. faithful. (As. 7)			✓	✓
Subgraph (As. 8)		✓		✓
Faithful. (As. 9)		✓		✓

Table 1: Assumptions required by each algorithm for soundness.

Corollary 1. *Let $\mathcal{D} = (\mathbb{V}, \mathbb{E})$ be the true FT-DFG, $\mathcal{A} \in \{\text{tsLDiffPC}^2, \text{tsDCIPC}\}$ and $\widehat{\mathcal{D}} = (\mathbb{V}, \widehat{\mathbb{E}}, \widehat{\mathbb{E}}^-)$ be the output of $\mathcal{A}$. Under the assumptions required by $\mathcal{A}$ in Table 1, and given perfect conditional equality information for all pairs of variables, $\text{Skel}(\mathcal{D}) = \text{Skel}(\widehat{\mathcal{D}})$ and $\widehat{\mathbb{E}} \subseteq \mathbb{E}$.*

Example 3. In Setting R, tsLDiffPC cannot orient the edge in the FT-DFG. However, tsLDiffPC² leverages orientations learned by tPC in the normal regime. Since the corresponding FT-DAG contains a collider with non-adjacent extremities, tPC correctly orients the edge, and this orientation is subsequently transferred to the difference graph. In Setting RB, tsLDiffPC² produces the same output as tsLDiffPC. Similarly, tsDCIPC combines residual-variance-based orientations with orientations obtained from tPC. Consequently, it correctly orients the edge in both settings. The outputs of both algorithms are shown in Figure 3 (middle) for Setting R and in Figure 3 (right) for Setting RB.

5 Difference graph discovery for detecting root causes

Up to this point, we have not directly addressed the detection of root causes. Once the FT-DFG has been estimated, root-cause detection follows naturally.

Under the assumptions of Table 1, every directed edge returned by an algorithm $\mathcal{A}$ corresponds to a true changed edge in the FT-DFG. Hence, any variable targeted by such an edge is a confirmed effect-defying root cause. In contrast, an undirected edge $X_t - Y_t$ only implies that at least one of X or Y is a root cause; we therefore refer to (X, Y) as a pair of *potential root causes*.

To formalize these recovery guarantees, define the set of certain root causes and the set of pairs of potential root causes as $\widehat{\mathcal{C}} := \{Y : \exists X_{t-\ell} \in \mathbb{V}, X_{t-\ell} \rightarrow Y_t \in \widehat{\mathcal{D}}\}$, and $\widehat{\mathcal{C}}_P := \{(X, Y) : X_t - Y_t \in \widehat{\mathcal{D}}\}$. Corollary 2 establishes the recovery guarantees for both sets.

Corollary 2. *Let $\mathcal{C}$ denote the set of effect-defying root causes, and let $\hat{\mathcal{C}}$ and $\hat{\mathcal{C}}_P$ be as defined above, for an algorithm $\mathcal{A} \in \{\text{tsLDiffPC}, \text{tsLDiffPC}^2, \text{tsDCI}, \text{tsDCIPC}\}$. Suppose that the assumptions required by $\mathcal{A}$, as specified in Table 1, are satisfied, and we are given perfect conditional equality information for all pairs of variables. Then: $\hat{\mathcal{C}} \subseteq \mathcal{C}$ and $\forall (X, Y) \in \hat{\mathcal{C}}_P, [X \in \mathcal{C}] \vee [Y \in \mathcal{C}]$. Furthermore among all the possible ways of selecting one variable from each pair in $\hat{\mathcal{C}}_P$, there exists a selection $\hat{\mathcal{C}}_S$ such that $\hat{\mathcal{C}} \cup \hat{\mathcal{C}}_S = \mathcal{C}$.*

The corollary implies that every detected root cause is a true root cause, i.e., $\hat{\mathcal{C}} \subseteq \mathcal{C}$. However, not all root causes are necessarily identified with certainty, since an undirected edge only indicates that only one of its endpoints is a root cause. The second item guarantees that all such unresolved root causes are contained in the set of potential root causes $\hat{\mathcal{C}}_P$.

Example 4. In both settings R and RB, $\mathcal{C} = \{Y\}$. On the left panel, tsLDiffPC only identifies *potential root causes*: $\hat{\mathcal{C}} = \emptyset$ and $\hat{\mathcal{C}}_P = \{(X, Y)\}$ meaning that X or Y is the root cause. In the middle panel tsLDiffPC² and tsDCIPC directly identifies the root cause, i.e. $\hat{\mathcal{C}}_P = \emptyset$ so $\hat{\mathcal{C}} = \{Y\} = \mathcal{C}$. For setting RB, all four methods correctly identify the root cause, i.e. $\hat{\mathcal{C}}_P = \emptyset$ and $\hat{\mathcal{C}} = \{Y\} = \mathcal{C}$ for tsLDiffPC, tsLDiffPC², and for tsDCIPC and tsDCI provided that $\sigma_Y^N = \sigma_Y^R$.

6 Experiments

In this section, we evaluate the proposed algorithms on both simulated and real-world datasets. We compare tsLDiffPC, tsLDiffPC², tsDCI, and tsDCIPC against two root cause analysis methods, namely MicroCause and RCD. We also consider naive extensions of existing graph discovery algorithms, namely tsMBGH and tsISCAN, derived from MBGH, iSCAN, and tPCUnion, which estimates a temporal graph separately in each regime and compares direct effects of the normal and anomalous regimes that are identified using the union graph.

For tsLDiffPC, tsLDiffPC², tsDCI, and tsDCIPC, the skeletons are estimated using a linear regression coefficient equality test across regimes. Edge orientations in tsDCI and tsDCIPC are obtained through an F -test for equality of residual variances across regimes. The tPC procedure used by tsLDiffPC² and tsDCIPC employs Fisher’s Z test for conditional independence testing. All other methods are run with their default hyperparameters. All statistical tests are performed at significance level $\alpha = 0.05$, and in all algorithms, ℓ_{max} is set to one. A sensitivity analysis with respect to the significance level is provided in Appendix.

6.1 Simulated data

We evaluate the algorithms on simulated multivariate time-series data generated from random finite-lag causal graphs. First, a random FT-DAG is sampled with a predefined number of time-series ($d \in \{3, 5, 7, 9\}$), edge probability 0.3 and $\ell_{max} = 1$. Edge coefficients are sampled uniformly from $[0.2, 0.8]$ with random

sign. To introduce anomalies, we consider three settings. In all cases, a target vertex is selected as an effect-defying root cause, and its incoming edge coefficients are modified while keeping the graph structure unchanged. In the first setting, only one incoming edge to the target vertex is modified; in the second setting, all incoming edges of the target vertex are modified; in the third setting, all incoming edges to the target vertex are modified, and we additionally impose that the target vertex has at least two parents. This last setting allows us to focus on cases where tsLDiffPC has sufficient information to orient edges.

Time series are then simulated recursively from the corresponding linear DT-DSCM with Gaussian noise. For each setting, we simulate 10 datasets. To evaluate the different algorithms, we use the F1-score defined with respect to the set of true root causes $\mathcal{C}$ and a set of predicted positives $\mathcal{P}$, *i.e.*, $F1(\mathcal{C}, \mathcal{P}) = \frac{2|\mathcal{C} \cap \mathcal{P}|}{|\mathcal{C}| + |\mathcal{P}|}$, where two choices are considered for the set of predicted positives $\mathcal{P}$: $\mathcal{P} = \hat{\mathcal{C}}$, where only confirmed root causes are considered; $\mathcal{P} = \hat{\mathcal{C}} \cup \hat{\mathcal{C}}_{\mathcal{P}}$ where both confirmed and potential root causes are considered.

Figure 4 shows that, when only confirmed root causes are considered, tsDCI and tsDCIPC achieve the best overall performance in the first and second settings. In contrast, tsLDiffPC often fails to orient the detected changed edges, although tsLDiffPC² improves its performance. In the third setting, tsLDiffPC and especially tsLDiffPC² become competitive with the DCI-based methods. This setting is particularly favorable to tsLDiffPC, as multiple changed parents are more likely to form unshielded colliders that facilitate edge orientation. When potential root causes are also taken into account, the performance of tsLDiffPC improves in all settings, confirming that its main limitation lies in the fact that it discovers a partially oriented graph. Among the remaining methods, RCD is the most competitive baseline, whereas tsMBGH, MicroCause, tPCUnion, and especially tsiSCAN generally obtain lower F1-scores.

6.2 IT monitoring

We evaluate the algorithms on IT monitoring data, namely, on the Ingestion dataset, introduced in [4]. The dataset was collected with a one-minute sampling rate and contains eight time-series measured over normal and anomalous operating periods. The time-series describe different activities of the ingestion pipeline, including message preprocessing (PMDB), message routing (MDB), metric extraction (CMB), database insertion (MB), updates of recent metric values (LMB), merging with check-message information (RTMB), insertion of historical status information (GSIB), and writing data to Elasticsearch (ESB). Each value is computed by multiplying the number of messages processed by the corresponding component during a 10-minute window by the average execution latency in that window, and then normalizing by the window length. A visualization of the time series is provided in Appendix.

The dataset contains 100 anomalous observations and 1000 normal observations preceding the anomalous period. It was highlighted in [4] that according to system experts two time-series were identified as root causes of the anomaly.




Fig. 4: Mean F1-score for nine methods as a function of the number of vertices across three settings.

lies. The first corresponds to PMDB; however, based on the expert knowledge discussed in [4], this time-series is not caused by any other observed time-series, making it unlikely to be an effect-defying root cause. The second corresponds to ESB, which is more plausibly an effect-defying root cause, as it was detected as such by EasyRCA using data and the available background knowledge.

On this dataset, all tsLDiffPC-based and tsDCI-based algorithms inferred the same graph, shown in Appendix. From this graph we can deduce that all these algorithms detected the effect-defying root cause ESB, along with three false positives. By contrast, tsMBGH did not detect any root cause, while tPCUnion and tsiSCAN identified all time-series as root causes (except RTMB for tsiSCAN). Among the other baseline methods, RCD identified LMB, MB and PMDB, whereas MicroCause detected CMB and PMDB. Overall, the tsLDiffPC-based and tsDCI-based methods are the only approaches that consistently recover ESB.

6.3 Intensive care monitoring

We illustrate the algorithms on a patient record from the MIMIC-IV Waveform Database [13], a publicly available critical care database of high-resolution physiological signals from ICU patients. Nine variables are retained: mean arte-

rial blood pressure (ABPm), mean pulmonary arterial pressure (PAPm), mean umbilical arterial pressure (UAPm), heart rate (HR), ST-segment deviation on leads V and III (ST-V, ST-III), central venous pressure (CVPm), respiratory rate (RR), and oxygen saturation (SpO2). Since the raw signals are irregularly sampled with missing values, all variables are first aligned onto a common time grid via linear interpolation. Regime change detection is performed on ST-V using the PELT algorithm [15] with an MBIC penalty, yielding a single changepoint. A symmetric window of 10 000 observations around this changepoint defines two temporal segments. For consistency with our notation, we refer to the earlier segment as the normal regime and to the later segment as the anomalous regime; these labels indicate temporal order only, not clinical status. A visualization of the time series is provided in Appendix.

The estimated FT-DFGs are shown in Appendix. The tsLDiffPC-based methods identify the same set of candidate root causes, namely RR, HR, PAPm, ABPm and CVPm. The tsDCI-based methods yield a similar result, but leave an ambiguity between PAPm and CVPm due to an unresolved edge orientation. The tsSCAN algorithm identifies RR as the sole root cause. In contrast, tP-Union detects all variables as root causes except SpO2. MicroCause identifies PAPm and CVPm as root causes, partially overlapping with the variables detected by the difference graph methods. And RCD identifies ST-V, which is not selected by any other method. Although no ground truth is available, there is an agreement among most algorithms on PAPm, CVP, and RR.

7 Conclusion

In this paper, we investigated how difference-graph discovery methods can be leveraged for root-cause analysis in time-series data. Our results demonstrate the potential of these methods for identifying structural changes associated with observed anomalies. The proposed approaches nevertheless rely on several assumptions regarding the underlying causal structure and data-generating processes. Moreover, their computational cost increases rapidly with both the number of variables and the maximum time lag considered, which may limit their applicability to large-scale systems.

Acknowledgment This work was supported by the CIPHOD project (ANR-23-CPJ1-0212-01).

References

1. Charles Assaad. *Contributions to causal reasoning and root cause analysis*. Habilitation à diriger des recherches, Sorbonne Université, January 2026.
2. Charles K. Assaad. Causal reasoning in difference graphs. In Mathias Drton and Biwei Huang, editors, *Proceedings of the Fourth Conference on Causal Learning and Reasoning*, Proceedings of Machine Learning Research. PMLR, 2025.

3. Charles K. Assaad, Emilie Devijver, and Eric Gaussier. Survey and evaluation of causal discovery methods for time series. *J. Artif. Int. Res.*, 73, May 2022.
4. Charles K. Assaad, Imad Ez-Zejjari, and Lei Zan. Root cause identification for collective anomalies in time series given an acyclic summary causal graph with loops. In Francisco Ruiz, Jennifer Dy, and Jan-Willem van de Meent, editors, *Proceedings of The 26th International Conference on Artificial Intelligence and Statistics*, volume 206 of *Proceedings of Machine Learning Research*, pages 8395–8404. PMLR, 25–27 Apr 2023.
5. K. Budhathoki, D. Janzing, P. Bloebaum, and H. Ng. Why did the distribution change? In *International Conference on Artificial Intelligence and Statistics*, pages 1666–1674. PMLR, 2021.
6. K. Budhathoki, L. Minorics, P. Blöbaum, and D. Janzing. Causal structure-based root cause analysis of outliers. In *International Conference on Machine Learning*, pages 2357–2369. PMLR, 2022.
7. Daria Bystrova and Emilie Devijver. Constraint-based difference graph discovery in a linear setting, 2026.
8. Daria Bystrova, Emilie Devijver, Vardan Manucharian, Julie Mondet, and Pascal Mossuz. Difference graph over two populations: Implicit difference inference algorithm. In *9th Causal Inference Workshop at UAI 2024*, 2024.
9. Gustau Camps-Valls, Miguel-Ángel Fernández-Torres, Kai-Hendrik Cohrs, Adrian Höhl, Andrea Castelletti, Aytac Pacal, Claire Robin, Francesco Martinuzzi, Ioannis Papoutsis, Ioannis Prapas, Jorge Pérez-Aracil, Katja Weigel, Maria Gonzalez-Calabuig, Markus Reichstein, Martin Rabel, Matteo Giuliani, Miguel D. Mahecha, Oana-Iuliana Popescu, Oscar J. Pellicer-Valero, Said Oualla, Sancho Salcedo-Sanz, Sebastian Sippel, Spyros Kondylatos, Tamara Happé, and Tristan Williams. Artificial intelligence for modeling and understanding extreme weather and climate events. *Nature Communications*, 16(1), 2025.
10. Tianyu Chen, Kevin Bello, Bryon Aragam, and Pradeep Ravikumar. iscan: Identifying causal mechanism shifts among nonlinear additive noise models. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 44671–44706. Curran Associates, Inc., 2023.
11. Alistair Connell and Chris Laing. Acute kidney injury. *Clinical medicine*, 15(6):581–584, 2015.
12. Azam Ikram, Sarthak Chakraborty, Subrata Mitra, Shiv Saini, Saurabh Bagchi, and Murat Kocaoglu. Root cause analysis of failures in microservices through causal discovery. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems*, volume 35, pages 31158–31170. Curran Associates, Inc., 2022.
13. Alistair E. W. Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Alon Shammout, Leo Lehman, Tom J. Pollard, Steven Horng, Leo Anthony Celi, and Roger G. Mark. MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, 10(1):1, 2023.
14. Kidney Disease: Improving Global Outcomes (KDIGO) Acute Kidney Injury Work Group. KDIGO Clinical Practice Guideline for Acute Kidney Injury. *Kidney International Supplements*, 2(1):1–138, 2012.
15. Rebecca Killick and Idris A. Eckley. **changepoint**: An R package for changepoint analysis. *Journal of Statistical Software*, 58(3):1–19, 2014.
16. Mingjie Li, Zeyan Li, Kanglin Yin, Xiaohui Nie, Wenchu Zhang, Kaixin Sui, and Dan Pei. Causal inference-based root cause analysis for online service systems with

- intervention recognition. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, KDD '22, page 3230–3240, New York, NY, USA, 2022. Association for Computing Machinery.
17. Nader A Mahmood, Fawad A Chaudry, Hamad Azam, M Imran Ali, and M Anees Khan. Frequency of hypoxic events in patients on a mechanical ventilator. *International journal of critical illness and injury science*, 3(2):124–129, 2013.
 18. Vineet Malik, Kevin Bello, Asish Ghoshal, and Jean Honorio. Identifying causal changes between linear structural equation models. In *Proceedings of the Fortieth Conference on Uncertainty in Artificial Intelligence*, UAI '24. JMLR.org, 2024.
 19. Christopher Meek. Causal inference and causal explanation with background knowledge. In *Conference on Uncertainty in Artificial Intelligence*, 1995.
 20. Yuan Meng, Shenglin Zhang, Yongqian Sun, Ruru Zhang, Zhilong Hu, Yiyin Zhang, Chenyang Jia, Zhaogang Wang, and Dan Pei. Localizing failure root causes in a microservice through causality inference. In *2020 IEEE/ACM 28th International Symposium on Quality of Service (IWQoS)*, pages 1–10, 2020.
 21. W. R. Orchard, N. Okati, S. Garrido Mejia, P. Blöbaum, and D. Janzing. Root cause analysis of outliers with missing structural knowledge. *Advances in Neural Information Processing Systems*, 38:161634–161671, 2026.
 22. Judea Pearl. *Causality: Models, Reasoning and Inference*. Cambridge University Press, USA, 2nd edition, 2009.
 23. Nicolas-Domenic Reiter, Jonas Wahl, Andreas Gerhardus, and Jakob Runge. Causal inference on process graphs: Causal structure and effect identification. *Bernoulli*, 32(2):1356 – 1382, 2026.
 24. Jakob Runge. Discovering contemporaneous and lagged causal relations in autocorrelated nonlinear time series datasets. In Jonas Peters and David Sontag, editors, *Proceedings of the 36th Conference on Uncertainty in Artificial Intelligence (UAI)*, volume 124 of *Proceedings of Machine Learning Research*, pages 1388–1397. PMLR, 03–06 Aug 2020.
 25. Jakob Runge, Sebastian Bathiany, Erik Bollt, Gustau Camps-Valls, Dim Coumou, Ethan Deyle, Clark Glymour, Marlene Kretschmer, Miguel D. Mahecha, Jordi Muñoz-Marí, Egbert H. van Nes, Jonas Peters, Rick Quax, Markus Reichstein, Marten Scheffer, Bernhard Schölkopf, Peter Spirtes, George Sugihara, Jie Sun, Kun Zhang, and Jakob Zscheischler. Inferring causation from time series in earth system sciences. *Nature Communications*, 10(2553), 2019.
 26. Daniela Schkoda and Dominik Janzing. Root cause analysis of outliers in unknown cyclic graphs. In *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2026. Poster.
 27. Peter Spirtes, Clark Glymour, and Richard Scheines. *Causation, Prediction, and Search*. MIT press, 2nd edition, 2000.
 28. Yuhao Wang, Chandler Squires, Anastasiya Belyaeva, and Caroline Uhler. Direct estimation of differences in causal graphs. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018.
 29. Lei Zan, Charles K. Assaad, Emilie Devijver, Eric Gaussier, and Ali Aït-Bachir. On the fly detection of root causes from observed data with application to it systems. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, CIKM '24, page 5062–5069, New York, NY, USA, 2024. Association for Computing Machinery.